



theme: gaia

\_class: lead

paginate: true

backgroundColor: #fff

[//]: # (Slide author: Fred Gras)

## **The Kubernetes Book** by Nigel Poulton *2021 Edition*

### **Chapter 4, Working w/ Pods**

**Feb 10, 2022** Austin OWASP Study Group

# What is Pod -- the atomic unit of scheduling

You cannot run a "container" directly in Kubernetes, needs to be wrapped in **Pod**

- Pods augment containers in the following ways
  - Labels and annotations
  - Restart policies
  - Probes (ie: startup, readiness, liveness, and more)
  - Affinity and anti-affinity rules
  - Termination controls
  - Resource sharing (requests and w/ limits)
- Pods can be scheduled
- Pods enable resource sharing

## Example of: \$ **kubectl explain pods --recursive**

```
KIND:      Pod
VERSION:   v1
```

### DESCRIPTION:

```
Pod is a collection of containers that can be run on a host.
This resource is created by clients and scheduled onto hosts.
```

### FIELDS:

```
  apiVersion    <string>
  kind           <string>
  metadata      <Object>
    annotations <map[string]string>
    labels      <map[string]string>
    name        <string>
    namespace   <string>
```

```
etc...
```

- notes: see example in two slides

## Example of: \$ `kubectl explain spec.restartPolicy`

KIND: Pod

VERSION: v1

FIELD: restartPolicy <string>

DESCRIPTION:

Restart policy for all containers within the pod.

One of **Always**, **OnFailure**, **Never**.

Default to **Always**.

More info: <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle/>

## Example of: labels, pods, affinities, policies, and dependencies

```
kind:      Pod
apiVersion: v1
metadata:   // this is the wrapper
  name:     hello-pod
  labels:
    zone:    prod
    version: v1
spec:
  containers: // this is a container (that is wrapped)
  - name: hello-ctr
    image: nigelpoulton/k8sbook:1.0
    ports:
    - containerPort: 8080
```

- this pod is a wrapper around one or more containers
- You will see this again many more times :)

## Pods enable **resource sharing**

- Shared filesystem
- Shared network stack (IP address , ports ...)
- Shared memory
- Shared volumes

note: ... will cover these details through the book (ie: not here :)

# Static vs controller Pods

Two ways to deploy pods

- Directly w/ a Pod manifest **Called static**
- Indirectly w/ a controller

Pods compared to cattle, when they die, they get replaced (ie: not a pet :)

## Single-container or multi-container Pods

(see later description)

# Deploying Pods

## steps as follows:

- Define it in a YAML manifest file
- Post the YAML to the API server
- The API server authenticates and authorizes the request
- The configuration (YAML) is validated
- The scheduler deploys the Pod to a healthy node with enough available resources
- The load kubelet monitors it

TODO: find Reference to confirm same steps



## the **Guts** of a Pod

a Pod is actually a special type of container called a **pause container**

pods are a collection of resources

these resources are actual Linux kernel namespaces, and have following

- **net namespace:** IP address, port range, routing tables...
- **pid namespace:** isolated process tree
- **UTS namespace:** Hostname
- **IPC namespace:** Unix domain sockets and shared memory

(not shown) Diagram showing a **private** pod network, do not need port mapping  
[similar image](#)

## **Atomic deployment of Pods -- Meaning all-or-nothing**

**Pod *lifecycle* -- pending then running, then terminate after succeeded state**

Short or long-lived Pods

**Pod *immutability* -- you can't modify them after deployed**

**Pods are ideal for scaling -- called *horizontal scaling***

## Summary about pods before moving on

- Pods are an atomic unit of scheduling
- single pods are simple, multi-container pods ideal (and fundamental) for tightly coupled workloads
- Pods are scheduled on physical nodes which do not span multiple nodes
- Pods are declared in a **manifest file** you post to the API server
- Pods deployed by higher-level controllers

note: leads into multi-container Pods

# Multi-container Pods -- each pod should have a clear defined responsibility

sometimes important to couple two or more functions

**co-locating (in the same Pod) allows containers to be designed for single responsibility such as:**

- **Sidecar** pattern -- the secondary task augments the main application container
- **Adapter** pattern -- the secondary takes main container output and reformats for external system
- **Ambassador** pattern -- the helper container **brokers** connectivity to external system
- **Init** pattern -- the special **init** container guarantees start/complete before run main application

## **Hands-on with Pods -- the second part of Chapter 4**

access the following repository for the book:

\$ git clone <https://github.com/nigelpoulton/TheK8sBook.git>

## Pod manifest files explained -- under pods (folder) called [pod.yml](#)

same example as before, add // comment

```
kind:      Pod      // object
apiVersion: v1      // normally <api-group>/<version>  v1 example StorageClass
metadata:   // where you attach things
  name:     hello-pod
  labels:   // default namespace
    zone:   prod
    version: v1
spec:      // spec will define the containers (desired state)
  containers: // just one -- this is single container pod
  - name:     hello-ctr
    image:    nigelpoulton/k8sbook:1.0
    ports:
    - containerPort: 8080
```

## Manifest files: Empathy as Code

Nigel mentioned *Nirmal Mehta* @ 2017 dockercon talk title [A Strong Belief, Loosely Held: Bringing Empathy to IT](#)

ie: the manifest(s) s/b excellent sources of documentation

## Deploying Pods from a manifest file -- **kubectl apply**

```
$ kubectl apply -f pod.yml
# // wait some time ...
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-pod	0/1	ContainerCreating	0	9s

# Introspecting running Pods -- **kubectl get**

- **-o wide** gives more columns
- **-o yaml** return full copy of the Pod from the cluster store

Output defined in two parts: **desired state** (.spec) and **observed state** (.status)

```
$ kubectl get pods hello-pod -o yaml
#... // yaml-snip
spec:                                // desired state
  containers:
  - image: nigelpoulton/k8sbook:1.0
    imagePullPolicy: IfNotPresent
    name: hello-ctr
#... // yaml-snip
status:                              // current observed state
  conditions:
  - lastProbeTime: null
    lastTransitionTime: "2022-02-08T18:21:51Z"
    status: "True"
    type: Initialized
```

- [put full listing in appendix](#) TODO



# kubectl describe

```
$ kubectl describe pods hello-pod
Name:          hello-pod
Namespace:     default
Start Time:    Mon, 09 Feb 2022 18:21:51 +0000
Labels:        version=v1
               zone=prod
Status:        Running
IP:            10.42.1.28
Containers:
  hello-ctr:
    Container ID:  containerd://<blah>
    Image:         nigelpoulton/k8sbook:1.0
    Port:          8080/TCP
# ... // snip
Conditions:
  Type           Status
  Initialized     True
  Ready           True
  ContainersReady True
# ... // snip
Events:
  Type    Reason      Age    Message
  ----    -
  Normal  Scheduled   5m30s  Successfully assigned ...
  Normal  Pulling     5m30s  Pulling image "nigelpoulton/k8sbook:1.0"
  Normal  Pulled      5m8s   Successfully pulled image ...
  Normal  Created     5m8s   Created container hello-ctr
  Normal  Started     5m8s   Started container hello-ctr
```

- [put full listing in appendix](#) TODO

## kubectl logs

```
$ kubectl logs <multipod>
# ... // yaml-snip
spec:
  containers:
    - name: app                      // first container
      image: nginx
      ports:
        - containerPort: 8080
    - name: syncer                  // second container
      image: k8s.gcr.io/gi-sync:v3.1.6
      volumeMounts:
        - name: html
# ... // etc
```

\$ kubectl logs multipod --container syncer // Example for specific container

## kubectl exec: running commands in Pods

execute commands within the pod

```
$ kubectl exec hello-pod -- ps aux
PID      USER     TIME    COMMAND
  -1      root      0.00    node ./app.js
  11      root      0.00    ps aux
```

get shell access

```
$ kubectl exec -it hello-pod -- sh // add --container if multi-container

// interactive in shell
# apk add curl
// then
# curl localhost:8080
```

## Pod **hostnames** -- command is case sensitive

container in Pod will inherit hostname from name of the pod  
valid DNS names (a-z 0-9, the minus-sign, and the period-sign)

## Pod **immutability**

```
$ kubectl edit <name>
```

Kubernetes will prevent the change

# Example multi-containers

**init container**

**sidecar container**

# The Chapter ending and Summary

## Cleanup

\$ `kubectl delete` pod <name>

## Chapter 4 Summary -- whats next ...

REF: eBook: TheUltimate Guide to Kubernetes Security

- mention in Austin-OWASP slack channel #kubernetes Jan27th (TODO URL)
- mention like the MARP-TEAM -- nice conversion from markdown to PDF

# **Appendix -- More detailed information**

**TODO**